

深圳大学考试答题纸

(以论文、报告等形式考核专用)

二〇____~二〇____ 学年度第____学期

课程编号	1500430001	课序号	02	课程名称	专业英语	主讲教师	杨景媛	评分	
学号	2024040042	姓名	汉雨	专业年级	2024 级计算机科学与技术创新班 02				

教师评语:

题目: **Image-Based Virtual Try-On: From Deep Learning Pipelines to Diffusion Models**

1. Abstract

Image-based virtual try-on is an economical and practical approach in the virtual try-on domain, aiming to synthesize realistic fitting results by combining person images with target garment images. It has attracted increasing attention in applications such as online shopping, fashion design, and digital content creation. In recent years, generative large-scale models, particularly diffusion models, have driven significant breakthroughs in this field due to their superior generative capability and training stability compared with traditional deep learning methods. However, a systematic analysis of image-based virtual try-on in the era of large generative models remains limited.

This paper presents a survey of image-based virtual try-on methods by organizing existing studies according to a standard three-stage pipeline: data preprocessing, garment warping, and try-on result generation. Representative approaches are reviewed and compared from both methodological and architectural perspectives. In addition, commonly used datasets, evaluation metrics, and loss functions are briefly summarized. Finally, current challenges and limitations of image-based virtual try-on are

discussed, and potential future research directions are outlined in the context of diffusion-based generative models.

2. Related Work

2.1 Background and Methodological Overview

With the rapid growth of e-commerce and digital content industries, virtual try-on has become an important research topic in computer vision. In the context of Artificial Intelligence Generated Content (AIGC), virtual try-on techniques

are not only widely adopted in online shopping scenarios to enhance user experience and decision-making efficiency but are also increasingly used in image and video content creation, enabling more personalized and editable garment synthesis and appearance manipulation.

From a methodological perspective, existing virtual try-on approaches can be broadly categorized into two paradigms: 3D virtual try-on methods and image-based virtual try-on methods. These two paradigms differ significantly in terms of data requirements, computational complexity, and practical applicability.

2.2 3D and Image-Based Virtual Try-On Methods

3D virtual try-on methods typically rely on accurate 3D human body reconstruction and physics-based cloth simulation, aiming to explicitly model the interaction between garments and the human body. One representative line of research focuses on physically based simulation, where garment deformation is modeled through cloth simulation combined with 3D body measurement and reconstruction techniques [2–5]. By explicitly modeling material properties, physical constraints, and lighting conditions, these methods offer high controllability and realism in garment geometry and appearance, making them particularly suitable for high-fidelity fashion design, virtual prototyping, and professional visualization applications.

However, despite their visual realism, physics-based 3D virtual try-on methods usually require specialized hardware, accurate body scanning, and complex processing pipelines, resulting in high computational cost and limited scalability. In addition, the need for expert knowledge and long simulation times further restricts their deployment in real-time or consumer-level scenarios. Another line of work seeks to

capture real garments using controlled hardware systems, such as robotic mannequins or multi-view capture setups, in order to achieve highly realistic visual results. While these approaches can preserve fine-grained garment details, they significantly sacrifice flexibility and incur substantial deployment costs, and are therefore only adopted in a small number of commercial systems, such as Fits.me [1].

In contrast, image-based virtual try-on methods have emerged as the dominant research direction in recent years due to their lower data acquisition cost and higher flexibility. These methods typically take 2D person images and garment images as input and can be trained and deployed under relatively limited computational resources, making them more suitable for large-scale e-commerce platforms and user-generated content scenarios. From a problem formulation perspective, image-based virtual try-on can be regarded as a conditional image generation task, where the model is required to synthesize realistic garment deformation, spatial alignment, and texture details while preserving the identity, pose, and background of the target person, often under weakly paired or even unpaired data settings.

2.3 From GAN-Based to Diffusion-Based Paradigms

Early image-based virtual try-on studies were primarily built upon generative adversarial networks (GANs). CAGAN [6] first introduced conditional generation for virtual try-on by compositing garment masks with person images but suffered from limited structural consistency and visual realism. Subsequently, VITON [7] proposed a representative three-stage pipeline consisting of try-on condition generation, explicit garment warping, and try-on image synthesis. This modular pipeline established a clear framework that guided a large body of subsequent research focusing on garment deformation modeling, image quality enhancement, and high-resolution synthesis [8–18].

In recent years, the research focus has gradually shifted from GAN-based frameworks to diffusion-based generative paradigms, following the success of diffusion models in general image synthesis tasks. Compared with GANs, diffusion models demonstrate superior training stability, stronger texture preservation, and increased robustness to complex poses and occlusions [19–31]. More importantly, the iterative denoising process of diffusion models enables them to jointly model garment deformation, alignment, and appearance synthesis within a unified generative framework, thereby reducing or even eliminating the need for explicit

garment warping modules and strictly separated pipeline stages. As a result, an increasing number of studies incorporate diffusion models into image-based virtual try-on systems and implicitly integrate garment transformation into end-to-end generation, marking a paradigm shift from modular “explicit warping + synthesis” pipelines toward fully implicit and unified generative models.

Although several surveys have summarized virtual try-on techniques from the perspectives of workflow and functionality [1,32–33], the rapid emergence of diffusion-based methods since 2023 has substantially reshaped the technical landscape. A comprehensive review that systematically analyzes image-based virtual try-on methods from the viewpoint of methodological paradigms and pipeline evolution remains limited. Therefore, it is necessary to revisit both classical approaches and recent diffusion-based advances to better understand their respective strengths and limitations, and to explore future research directions for image-based virtual try-on in the era of large generative models

3. Key studies

3.1 Early GAN-based Virtual Try-On: CAGAN

Jetchev and Bergmann proposed CAGAN (Conditional Analogy Generative Adversarial Network) as one of the earliest attempts to address image-based virtual try-on from a generative modeling perspective [6]. Instead of treating virtual try-on as a geometric transformation or explicit garment fitting problem, CAGAN formulates it as a conditional image generation task, in which the appearance of a target garment is transferred onto a person image while the person’s identity, body structure, and pose are preserved as much as possible.

Methodologically, CAGAN is built upon a generative adversarial network (GAN) framework that learns analogy relationships between paired images. Given a person image wearing an original garment and a separate target garment image, the model predicts both an appearance transformation and an implicit blending mask. This mask controls how the generated garment appearance is composited with the original person image, enabling the network to selectively replace clothing regions while retaining non-garment areas such as the face and background. Owing to this end-to-end design, CAGAN is able to perform clothing swapping using only 2D visual information, without relying on explicit 3D human models, physical cloth simulation, or

handcrafted geometric constraints.

A key advantage of CAGAN lies in its conceptual simplicity and unified training strategy. By implicitly learning blending masks during adversarial training, the method avoids the need for supervised clothing segmentation, pose estimation, or complex preprocessing pipelines, thereby significantly lowering the entry barrier for image-based virtual try-on research. More importantly, CAGAN establishes a foundational perspective by framing virtual try-on as a conditional image synthesis problem, which directly inspired a large body of subsequent work based on deep generative models. Nevertheless, CAGAN also exhibits several inherent limitations that restrict its practical applicability. Since the quality of the final try-on result strongly depends on the predicted blending mask, the generated images often suffer from visible artifacts around garment boundaries, especially when clothing contours are complex. In addition, CAGAN does not explicitly model geometric correspondence between the garment and the human body, nor does it account for garment deformation under different poses. As a result, the method struggles to handle challenging scenarios involving large pose variations, significant shape differences, or fine-grained texture details. These shortcomings highlight the necessity of more structured designs and motivate later pipeline-based approaches that explicitly decompose virtual try-on into multiple stages, such as garment alignment, warping, and image synthesis.

3.2 Pipeline-based Virtual Try-On: VITON

Han et al. proposed VITON as a representative pipeline-based approach for image-based virtual try-on, aiming to generate photo-realistic try-on images using only 2D RGB inputs without relying on 3D information [7]. To overcome the limitations of prior GAN-based methods in modeling geometric deformation and fine-grained details, VITON reformulates virtual try-on as a structured, coarse-to-fine generation process.

Specifically, the method decomposes the task into three stages: clothing-agnostic person representation generation, coarse try-on synthesis with clothing mask prediction, and refinement through garment warping and compositing. This explicit decomposition establishes a clear pipeline that separates pose and identity preservation, garment deformation, and visual detail synthesis.




Figure 1. Overview of the VITON framework, illustrating the clothing-agnostic person representation, coarse image synthesis, explicit garment warping, and refinement stages. The figure is reproduced from Han et al. (CVPR 2018).

A key contribution of VITON is the introduction of a clothing-agnostic person representation, which serves as a conditioning signal for generation. As described in Section 3.1 of the original paper, this representation encodes pose information using keypoint heatmaps, body shape information derived from human parsing, and identity-related regions such as face and hair. By removing clothing appearance from the conditioning input, the model is guided to focus on transferring the target garment while preserving the original person’s pose, body structure, and identity.

Conditioned on this representation and the target clothing image, VITON employs a multi-task encoder-decoder generator to produce a coarse synthesized image and a corresponding clothing mask. As shown in Figure 1 of the paper, the generator follows a U-Net-style architecture with skip connections and is trained using a combination of perceptual loss and L1 loss. While this coarse stage successfully preserves global structure and pose consistency, the authors observe that fine-grained garment details, such as logos and textures, are often missing in the generated results.

To compensate for this limitation, VITON introduces a refinement stage that explicitly models garment deformation and appearance preservation. The target clothing image is first warped using a thin-plate spline (TPS) transformation estimated via shape context matching between the clothing mask and the predicted garment region. The warped gar-

ment is then composited with the coarse output using a refinement network that predicts a composition mask, determining how much information should be taken from the warped garment versus the coarse synthesis. This design allows detailed visual patterns from the product image to be preserved while maintaining spatial consistency with the person's pose and body shape.

Experimental results on the newly collected Zalando dataset demonstrate that VITON produces more realistic and visually appealing try-on results than prior GAN-based approaches, including CAGAN. Both quantitative evaluations and user studies reported in the paper indicate that explicitly modeling person representation, garment warping, and refinement within a structured pipeline significantly improves virtual try-on quality. At the same time, the reliance on multiple hand-designed stages and explicit warping operations highlights the complexity of pipeline-based methods, which later work seeks to simplify or replace with more unified generative frameworks.

3.3 Warping-aware Pipeline Refinement: CP-VTON

Wang et al. proposed CP-VTON as an important extension of the pipeline-based virtual try-on framework, aiming to better preserve garment characteristics while maintaining alignment with the target person image [8]. Building upon the overall structure introduced by VITON, CP-VTON explicitly addresses the limitations of hand-designed geometric transformations and detail loss in previous pipeline-based methods.

The authors observe that in VITON, the thin-plate spline (TPS) transformation estimated by shape context matching is not learnable and often leads to inaccurate alignment when the target garment and the human body exhibit large shape discrepancies. To overcome this issue, CP-VTON reformulates garment warping as a learnable geometric matching problem and proposes a two-stage pipeline consisting of a Geometric Matching Module (GMM) and a Try-On Module.

The Geometric Matching Module is designed to estimate a TPS transformation in a fully learnable manner. As described in Section 3.1 of the original paper, GMM takes the clothing-agnostic person representation and the in-shop garment image as input, and learns dense correspondence features through convolutional layers. A regression network then predicts the parameters of the TPS transformation, enabling the garment to be spatially aligned with the target person in an end-to-end fashion. Compared with the

non-learnable warping strategy used in VITON, this approach significantly improves alignment accuracy and robustness to shape variation.




Figure 2. Overview of the CP-VTON framework. The framework consists of a Geometric Matching Module (GMM) for garment alignment and a Try-On Module for image synthesis. Reproduced from Wang et al. (ECCV 2018).

After geometric alignment, CP-VTON employs a Try-On Module to synthesize the final try-on image while preserving fine-grained garment details. This module follows an encoder-decoder architecture and predicts a composition mask that adaptively fuses the warped garment image with the rendered person representation. By explicitly learning how much information should be taken from the warped clothing versus the synthesized image at each pixel, CP-VTON effectively retains characteristic patterns such as logos, textures, and color distributions, which are often lost in coarse synthesis stages. This pixel-wise composition strategy plays a critical role in balancing visual realism and structural consistency in the final output.

Experimental results reported in the original study demonstrate that CP-VTON produces more visually faithful try-on images than VITON, particularly in terms of garment alignment accuracy and texture preservation. The improvements highlight the effectiveness of explicitly modeling geometric correspondence and composition in pipeline-based designs. However, despite these advances, CP-VTON still relies on a multi-stage pipeline with separately trained modules and explicit warping operations. Such structural complexity increases system design and training overhead, and also makes error propagation across stages difficult to control. These limitations motivate subsequent research toward more unified and implicit generative frameworks for virtual try-on, which aim to reduce manual module design while improving robustness.

3.4 Diffusion-based Virtual Try-On: TryOnDiffusion

Recent advances in diffusion models have significantly influenced image generation research, motivating their application to image-based virtual try-on. Zhu *et al.* proposed TryOnDiffusion, which is the first work to systematically introduce diffusion models into virtual try-on and to reconsider the traditional pipeline-based formulation from a unified generative perspective [19].

Conventional image-based virtual try-on methods typically decompose the task into explicit stages, such as garment warping followed by image synthesis and refinement. As discussed by the authors, this modular design often suffers from error accumulation, especially when handling large pose variations, severe occlusions, or complex garment shapes. In contrast, TryOnDiffusion reformulates virtual try-on as a conditional diffusion process, where garment deformation and appearance synthesis are jointly modeled rather than handled by separate, hand-designed modules.

A key technical contribution of TryOnDiffusion is the proposed Parallel-UNet architecture, which consists of two interacting UNets connected through cross-attention mechanisms. Instead of explicitly estimating pixel-wise displacement fields or warping transformations, the method performs garment alignment implicitly in the latent feature space. Cross-attention layers allow garment features to dynamically attend to person image features at multiple resolutions, enabling flexible and robust correspondence modeling under challenging pose and shape variations.

To support high-resolution image synthesis, TryOnDiffusion adopts a cascaded diffusion strategy. A base diffusion model first generates a coarse try-on result at low resolution, which is subsequently refined by higher-resolution diffusion models. This design follows the general principle of cascaded diffusion and allows the method to generate photo-realistic try-on images with detailed textures and fine-grained garment patterns.

Experimental results reported in the paper show that TryOnDiffusion outperforms previous GAN-based and pipeline-based virtual try-on methods in both quantitative evaluations and user studies, particularly in scenarios involving complex poses and body shapes. Nevertheless, the method still depends on preprocessing steps such as human parsing and pose estimation, and its performance may be affected by inaccuracies in these inputs. Despite this limitation, TryOnDiffusion represents a paradigm shift in image-based virtual try-on, marking a transition from explicit, stage-wise pipelines toward unified, diffusion-driven gener-

ation frameworks.

Subsequent studies, such as StableVITON [20], further explore how structural and semantic constraints can be incorporated into diffusion-based virtual try-on, indicating an ongoing effort to stabilize and improve this emerging paradigm.

4. Discussion

4.1 Evolution of Image-Based Virtual Try-On

Image-based virtual try-on has undergone a clear methodological evolution, progressing from early GAN-based image synthesis to pipeline-based frameworks and, more recently, diffusion-driven unified generative models. The analysis of representative studies in this work illustrates how each generation of methods incrementally addresses specific limitations of prior approaches while simultaneously reshaping the overall problem formulation.

Early GAN-based methods, such as CAGAN, demonstrate the feasibility of modeling virtual try-on as a conditional image generation task. By directly synthesizing try-on results from paired inputs, these approaches establish a conceptual foundation for image-based virtual try-on. However, due to their limited geometric reasoning capability, they often suffer from weak spatial consistency, inaccurate garment boundaries, and insufficient preservation of fine-grained texture details, particularly under pose variations.

Pipeline-based approaches, represented by VITON and CP-VTON, significantly improve visual quality and controllability by explicitly decomposing the task into human representation, garment warping, and result synthesis. This modular design enables more accurate garment alignment and deformation modeling, leading to more realistic results. Nevertheless, the introduction of multiple stages also increases system complexity, causes error accumulation across intermediate modules, and reduces robustness in challenging scenarios involving occlusions or extreme poses.

More recently, diffusion-based methods, exemplified by TryOnDiffusion, represent a paradigm shift by reformulating virtual try-on as a unified conditional generation process. By implicitly modeling garment deformation and appearance synthesis within a diffusion framework, these methods reduce dependence on hand-designed warping modules and explicit stage-wise coordination. This transition indicates that increasingly powerful generative models are capable of absorbing traditional pipeline components into end-to-end

frameworks, resulting in improved robustness and visual realism.

4.2 Remaining Challenges and Future Directions

Despite the substantial progress achieved by recent methods, several challenges remain unresolved. Current image-based virtual try-on systems still depend heavily on accurate pose estimation and semantic parsing, making their performance sensitive to errors in upstream human analysis. In addition, the diversity of existing datasets is limited in terms of body shapes, clothing styles, viewpoints, and real-world environments, which restricts generalization to unconstrained user scenarios.

Another important limitation lies in evaluation methodology. Commonly used metrics, such as SSIM, FID, and LPIPS, focus primarily on pixel-level or feature-level similarity and do not fully capture perceptual realism or user satisfaction, especially in weakly paired or unpaired settings. As a result, quantitative evaluations may not always align with human judgment of try-on quality. Furthermore, diffusion-based approaches introduce new practical challenges related to high computational cost, memory consumption, and inference latency, which hinder large-scale deployment and real-time applications.

Future research may benefit from richer and more robust human–garment representations, more diverse and fine-grained datasets, and evaluation criteria better aligned with human perception. In addition, exploring more efficient or hybrid generative frameworks could help balance generation quality with computational efficiency. Addressing these challenges is essential for advancing image-based virtual try-on from controlled research settings toward practical and user-oriented applications.

5. Conclusion

5.1 Summary of Core Contributions

This paper reviews recent advances in image-based virtual try-on, with a particular focus on the methodological evolution from early GAN-based approaches to structured pipeline designs and, more recently, diffusion-based generative frameworks. Rather than treating these methods in isolation, this survey organizes existing work along a coherent evolutionary trajectory, highlighting how different design paradigms emerge in response to the limitations of preceding approaches.

By analyzing representative studies, including CAGAN, VITON, CP-VTON, and TryOnDiffusion, we illustrate how

the field has progressively transitioned from direct image synthesis toward more unified and expressive generative models. In particular, this evolution reflects a shift from explicitly decomposed, stage-wise pipelines to end-to-end frameworks that implicitly model garment deformation, alignment, and appearance synthesis. Such a progression underscores the increasing capability of modern generative models to handle complex visual transformations while maintaining person identity and visual realism.

5.2 Broader Significance and Outlook

The transition from explicit, stage-wise pipelines to diffusion-driven generation represents a key contribution to the development of image-based virtual try-on. By shifting from modular designs toward unified generative frameworks, diffusion-based methods demonstrate improved robustness, visual realism, and flexibility, as garment transformations and appearance synthesis are implicitly modeled within an end-to-end process. This design paradigm reduces reliance on hand-designed warping operations and complex module coordination, thereby simplifying system architecture while enhancing overall generation quality.

Despite these advances, several limitations remain. Current approaches still depend heavily on accurate human pose estimation and semantic parsing, making them sensitive to errors in upstream human analysis. Moreover, their performance is constrained by limited dataset diversity and imperfect evaluation metrics, which may fail to reflect perceptual realism or user experience. In addition, the high computational cost and inference latency of diffusion models pose significant challenges for large-scale deployment and real-time applications.

Future research may benefit from richer and more robust human–garment representations, more diverse and perceptually aligned datasets and evaluation protocols, and more efficient generative architectures that better balance quality and efficiency. Addressing these challenges is essential for advancing image-based virtual try-on from controlled research settings toward practical, user-oriented applications, and for enabling broader adoption in real-world e-commerce and interactive systems.

References

- [1]. SONG D, ZHANG X P, ZHOU J, et al. Image-based virtual try-on: a survey[J]. arXiv:2311.04811, 2023.
- [2]. XU W W, UMENTANI N, CHAO Q W, et al. Sensitivity optimized rigging for example-based real-time clothing synthesis[J]. ACM Transactions on Graphics, 2014, 33(4): 107.
- [3]. WANG L Y, LI H L, XIAO Q J, et al. Automatic pose and wrinkle transfer for aesthetic garment display[J]. Computer Aided Geometric Design, 2021, 89: 102020.
- [4]. WU N N, CHAO Q W, CHEN Y Z, et al. AgentDress: real-time clothing synthesis for virtual agents using plausible deformations[J]. IEEE Transactions on Visualization and Computer Graphics, 2021, 27(11): 4107-4118.
- [5]. PAN X Y, MAI J M, JIANG X W, et al. Predicting loose-fitting garment deformations using bone-driven motion networks[C]//Proceedings of the ACM SIGGRAPH 2022 Conference. New York: ACM, 2022: 11.
- [6]. JETCHEV N, BERGMANN U. The conditional analogy GAN: swapping fashion articles on people images[C]//Proceedings of the 2017 IEEE International Conference on Computer Vision Workshops. Piscataway: IEEE, 2017: 2287-2292.
- [7]. HAN X T, WU Z X, WU Z, et al. VITON: an image-based virtual try-on network[C]//Proceedings of the 2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2018: 7543-7552.
- [8]. WANG B C, ZHENG H B, LIANG X D, et al. Toward characteristic-preserving image-based virtual try-on network[C]//Proceedings of the 15th European Conference on Computer Vision. Cham: Springer, 2018: 607-623.
- [9]. YU R Y, WANG X Q, XIE X H. VTNFP: an image-based virtual try-on network with body and clothing feature preservation[C]//Proceedings of the 2019 IEEE/CVF International Conference on Computer Vision. Piscataway: IEEE, 2019: 10510-10519.
- [10]. ISSENHUTH T, MARY J, CALAUZÈNES C, et al. End-to-end learning of geometric deformations of feature maps for virtual try-on[J]. arXiv:1906.01347, 2019.
- [11]. MINAR M R, TUAN T T, AHN H, et al. CP-VTON+: clothing shape and texture preserving image-based virtual try-on[C]//Proceedings of the 2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2020: 10-14.
- [12]. LEE H J, LEE R, KANG M, et al. LA-VITON: a network for looking-attractive virtual try-on[C]//Proceedings of the 2019 IEEE/CVF International Conference on Computer Vision Workshops. Piscataway: IEEE, 2019: 3129-3132.
- [13]. YANG H, ZHANG R M, GUO X B, et al. Towards photo-realistic virtual try-on by adaptively generating and preserving image content[C]//Proceedings of the 2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2020: 7847-7856.
- [14]. GE C, SONG Y, GE Y, et al. Disentangled cycle consistency for highly-realistic virtual try-on[C]//Proceedings of the 2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2021: 16928-16937.
- [15]. REN B, TANG H, MENG F Y, et al. Cloth interactive transformer for virtual try-on[J]. arXiv:2104.05519, 2021.
- [16]. CHOI S, PARK S, LEE M, et al. VITON-HD: high-resolution virtual try-on via misalignment-aware normalization[C]//Proceedings of the 2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2021: 14126-14135.
- [17]. LI K D, CHONG M J, ZHANG J, et al. Toward accurate and realistic outfits visualization with attention to details[C]//Proceedings of the 2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2021: 15541-15550.
- [18]. HAN X T, HUANG W L, HU X J, et al. ClothFlow: a flow-based model for clothed person generation[C]//Proceedings of the 2019 IEEE/CVF International Conference on Computer Vision. Piscataway: IEEE, 2019: 10470-10479.
- [19]. HU L Y, YANG D W, ZHU T, et al. TryOnDiffusion: a tale of two UNets[C]//Proceedings of the 2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2023: 4606-4615.
- [20]. KIM J, GU G, PARK M, et al. StableVITON: learning semantic correspondence with latent diffusion model

- for virtual try-on[C]//Proceedings of the 2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2024: 8176-8185.
- [21].XU Y H, GU T, CHEN W F, et al. OOTDiffusion: outfitting fusion based latent diffusion for controllable virtual try-on[J]. arXiv:2403.01779, 2024.
- [22].CHOI Y, KWAK S, LEE K, et al. Improving diffusion models for authentic virtual try-on in the wild[J]. arXiv:2403.05139, 2024.
- [23].YANG X, DING C X, HONG Z B, et al. Texture-preserving diffusion models for high-fidelity virtual try-on[C]//Proceedings of the 2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2024: 7017-7026.
- [24].WANG H Y, ZHANG Z L, DI D L, et al. MV-VTON: multi-view virtual try-on with diffusion models[J]. arXiv:2404.17364, 2024.
- [25].ZHU L Y, LI Y W, LIU N, et al. M&M VTO: multi-garment virtual try-on and editing[C]//Proceedings of the 2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2024: 1346-1356.
- [26].SHEN F, JIANG X, HE X, et al. IMAGDressing-v1: customizable virtual dressing[J]. arXiv:2407.12705, 2024.
- [27].CHONG Z, DONG X, LI H X, et al. CatVTON: concatenation is all you need for virtual try-on with diffusion models[J]. arXiv:2407.15886, 2024.
- [28].ZENG J H, SONG D, NIE W Z, et al. CAT-DM: controllable accelerated virtual try-on with diffusion model[C]//Proceedings of the 2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2024: 8372-8382.
- [29].LIN A R, ZHAO N X, NING S L, et al. FashionTex: controllable virtual try-on with text and texture[J]. arXiv:2305.04451, 2023.
- [30].XING J Z, XU C, QIAN Y J, et al. TryOn-adapter: efficient fine-grained clothing identity adaptation for high-fidelity virtual try-on[J]. arXiv:2404.00878, 2024.
- [31].LI H. Animate anyone: consistent and controllable image-to-video synthesis for character animation[C]//Proceedings of the 2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition. Piscataway: IEEE, 2024: 8153-8163.